



2024 Autumn Intro-to-Machine Learning Homework 1

Presenter: TA Wei-Hsiang (Sean)

Release Date: 2024/09/24 15:00

Homework 1

- Deadline: 23:59, Oct. 8th (Tue), 2024
- **Coding (60%)**: Implement linear regression by **only** using *numpy*.
 - Submit your code in executable python files (.py).
 - Report the outcome and parameters by screenshots to the questions.
- **Handwritten Questions (40%)**: Answer questions about linear regression.
 - Answer the questions in the report.
 - You must use the template and in digital-typed (no handwritten scan)
 - In English

Links

- [Questions and Report template](#)
- [Sample code / Dataset](#)

Coding Environment


- Recommendation: Python 3.9 or higher
- Tips
 - We recommend you to use **virtual environments** when implementing your homework assignments.
 - Here are some popular virtual environment management tools
 - [Poetry](#)
 - [Conda](#)
 - [Virtualenv](#)

Numpy

- High efficient vector and matrix operations
- Numpy Tutorial: [Link](#)

element-wise
multiply


```
a = np.array([1, 2, 3])  
b = np.array([4, 5, 6])  
for i in range(a.shape[0]):  
    a[i] *= b[i]  
print(a)  
# a = [ 4 10 18]
```



```
a = np.array([1, 2, 3])  
b = np.array([4, 5, 6])  
a *= b  
print(a)  
# a = [ 4 10 18]
```

square root

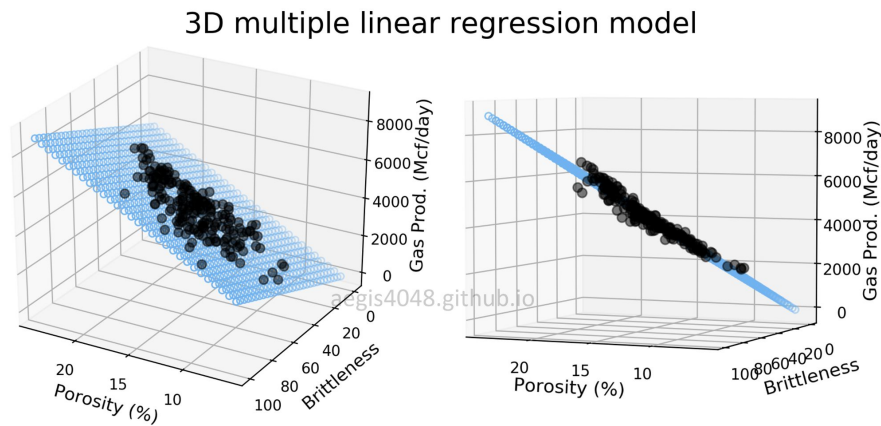
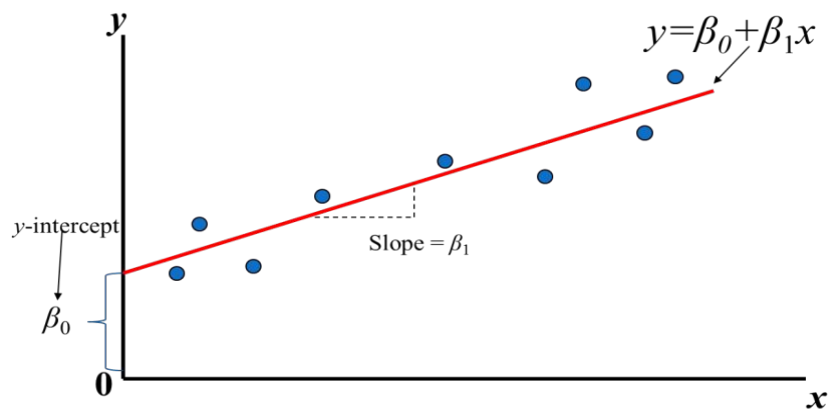
```
import math  
a = np.array([1, 4, 9])  
for i in range(a.shape[0]):  
    a[i] = math.sqrt(a[i])  
print(a)  
# a = [1 2 3]
```



```
a = np.array([1, 4, 9])  
a = np.sqrt(a)  
print(a)  
# a = [1 2 3]
```

Linear Regression

- Find the the slope (weights) and the intercept of given data



How to find β_0 and β_1 ?



- Implement the closed-form solution (*Question 1-1*)

$$\hat{\beta} = (X^T.X)^{-1}X^T.Y$$

- How about a huge dataset?
 - high dimensional data
 - huge amount of data


How to find β_0 and β_1 ?

- Gradient Descent (Question 1-2 ~ Question 1-4)



Gradient Descent

- x-axis and y-axis: the value of **weights**
- z-axis: the **value of loss** of the corresponding weights
- Goal: Find the weights that **minimize** the loss value



Dataset and Environment

- Student Performance Dataset
- Features
 - Hour Studied
 - Previous Score
 - Sleep Hours
 - Sample Question Papers Practiced
- Target
 - Performance Index (higher means better performance)
- Required packages: ``numpy``, ``pandas``, ``matplotlib``, ``loguru``, ``flake8``, ``pytest``

Linear Regression – Closed-form Solution (10%)

- Requirements
 - Implement Linear Regression by **closed-form** solution.
- Grading Criteria
 - (10%) Show the weights and intercepts of your linear model.
- Tips
 - There is only one answer.
 - You can check your answer by yourself using third-party libraries (such as scikit-learn).

Linear Regression – Gradient Descent (40%)

- Requirements

- Update your weights and intercept by using **gradient descent**
 - you can implement mini-batch gradient descent or stochastic gradient descent if you want.
- Use MSE (Mean Square Error) as your loss function.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Tune the **learning rate** and **epoch** hyper-parameters (and **batch size** if you implement mini-batch gradient descent) to make your testing MSE loss as close as the closed-form solution.

Linear Regression – Gradient Descent (40%)

- Grading Criteria

- (10%) Show the hyperparameters of your setting. Also, show the weights and intercepts of your fitted model.
- (10%) Plot the learning curve. (x-axis=epoch, y-axis=training loss)
- (20%) Show your MSE.cf, MSE.gd, and error rate between your closed-form solution and the gradient descent solution.
 - error rate: $|(\text{gradient_descent_loss} - \text{closed_form_loss}) / \text{closed_form_loss} | * 100$

Points	error rate, given both of the MSE of close-form and MSE of gradient descent < 20
20	< 0.5%
15	< 1%
10	< 3%
5	< 5%
0	>= 5%

Linear Regression – Gradient Descent

- Tips
 - Finding suitable hyper-parameters may cost you some time. Be patient!

Code Output

- Do not modify the main function architecture heavily.
- Your code output will look like this

```
2024-03-14 23:13:57.944 | INFO | __main__:main:77 - LR_CF.weights=
```

```
2024-03-14 23:54:18.052 | INFO | __main__:main:84 - LR_GD.weights=
```

```
2024-03-14 23:54:18.055 | INFO | __main__:main:93 - Prediction difference:
```

```
2024-03-14 23:54:18.055 | INFO | __main__:main:98 - mse_cf=, mse_gd=. Difference: 0.027%
```

1. [PEP8](#)
2. [Google Python Style](#)

Additional Requirements (10%)

Code Check and Verification: **Lint** the code and show the **PyTest** results (10%)

- Code linting: `$ flake8 main.py`
 - **-2pt** per warning / error
- Run PyTest: `$ pytest ./test_main.py -s`
 - **-5pt** per failed case

```
./main.py:103:1: W391 blank line at end of file
1      W391 blank line at end of file
```

```
===== test session starts =====
platform linux -- Python 3.9.5, pytest-8.0.2, pluggy-1.4.0
rootdir: /
collected 2 items

test_main.py 2024-03-16 11:52:21.189 | INFO | test_main:test_regression_cf:27 - model.weights=array([[3.]])
2024-03-16 11:52:21.190 | INFO | main:fit:57 - EPOCH 0, loss=3147.416663702691
2024-03-16 11:52:21.644 | INFO | main:fit:57 - EPOCH 10000, loss=0.29281584845965486
2024-03-16 11:52:22.094 | INFO | main:fit:57 - EPOCH 20000, loss=0.00536096424057785
2024-03-16 11:52:22.544 | INFO | main:fit:57 - EPOCH 30000, loss=9.815021195041223e-05
2024-03-16 11:52:22.998 | INFO | main:fit:57 - EPOCH 40000, loss=1.7969648133316264e-06
2024-03-16 11:52:23.450 | INFO | main:fit:57 - EPOCH 50000, loss=3.2899394472691304e-08
2024-03-16 11:52:23.905 | INFO | main:fit:57 - EPOCH 60000, loss=6.023324157052075e-10
2024-03-16 11:52:24.363 | INFO | test_main:test_regression_gd:39 - model.weights=array([3.]), model.intercept=3.9999966785390386
.

===== 2 passed in 4.37s =====
```


Handwritten Questions (40%)

2-1 (10%) How does the presence of outliers affect the performance of a linear regression model? How should outliers be handled? List at least two methods.

2-2 (15%) How do different values of learning rate (too large, too small...) affect the convergence of optimization? Please explain in detail.

2-3 (15%) What is the prior, likelihood, and posterior in Bayesian linear regression. (Explain the concept in detail rather than writing out the mathematical formula.) **What is the difference between Maximum Likelihood Estimation (MLE) and Maximum A Posteriori Estimation (MAP)?** (Analyze the assumptions and the results.)

Report

- Please follow the report template format. (-5pts if not use the template)
- [Link](#)

Submission

- Compress your code and report into a **.zip file** and submit it to E3.
- Report should be written in English. (-5 pts if not English)
- <STUDENT ID>_HW1.zip
 - main.py
 - setup.cfg
 - test_main.py
 - <STUDENT ID>_HW1.pdf (NO .doc, .docx or others format)
- Don't put the data (e.g. train.csv / test.csv) into submission file

Other rules

- **Late Policy**: A penalty of **20 points** per additional late day. (-20pt / delayed.day)
 - For example, If you get 90 points but delay for two days, your will get only 50 points!
- **No Plagiarism**: You should complete the assignment by yourself. Students engaged in plagiarism will be penalized heavily. Super serious penalty.
 - e.g. -100pt for the assignment or failed this course, etc
 - Report to academic integrity office

AI-Assistant

- Not recommended but no forbidden
- Copy-and-Paste answers from the AI-Assiant will be seen as Plagiarism
 - However, you can have your own answer first then rephrase it by AI-Assiant.
- Some questions might be parts of final exam, make sure you understand the concept



FAQs

- Why can't my gradient descent model converge?
 - Make sure you calculate the gradients correctly.
 - Use smaller learning rate.
- Can I use deep learning frameworks such as TensorFlow, PyTorch or other library such as math?
 - **No!** In HW1, you are request using **only Numpy** to implement linear regression and gradien descent. You can use matplotlib to plot the results.
- If you have other questions, ask on **E3 forum** first! We will reply as soon as possible.

Have Fun!

